

Lab 7: ELF-introduction

This lab **may** be done in pairs.

Lab Goals

- Extracting useful information from files in ELF.
- Fixing files using this information: reverse engineering.

In the following labs, you will learn to handle object and executable files. We will begin by learning just some of the basics of ELF, together with applications you can already use at this level - editing binary files and writing software patches. Then, we will continue our study of ELF files, by beginning to parse the structures of ELF files, and to use them for various purposes. In particular, we will access the data in the section header table and in the symbol table.

Methodology

- [Get to know the ELF.](#)
- Learn how to use the *readelf* utility. By using *readelf* you can get, in a human readable format, all the ELF structural information.
- Experience basic ELF manipulation.

Recommended Operating Procedure

This advice is relevant for all tasks. Note that while at some point you will no longer be using *hexedit* to process the file and *readelf* to get the information, nevertheless in some cases you may still want to use these tools for debugging purposes. In order to take advantage of these tools and make your tasks easier, you should:

- Support debugging messages: in particular the offsets of the various items, as you discover them from the headers.
- Use *hexedit* and *readelf* to compare the information you are looking for, especially if you run into unknown problems: *hexedit* is great if you know the exact location of the item you are looking for.
- Note that while the object files you will be processing will be linked using *ld*, and will, in most cases, use direct system calls in order to make the ELF file simpler, there is no reason why the programs you write need use this interface. You are allowed to use the standard library when building your own C programs.
- In order to preserve your sanity, even if the code you MANIPULATE may be without `stdlib`, we advise that for your OWN CODE you DO use the C standard library! (Yes, this is repeated twice, so that you notice it!)
- In order to keep sane in the following labs as well, **understand** what you are doing and **keep track** of that and of your code, as you will be using them in future labs.

All the executable files we will work with in this session are 32-Bit ELF binaries.
Compile your code accordingly.

Lab 7 Tasks

Deliverables

You should read and understand the reading material, and do task 0 before attending the lab. To be eligible for a full grade, you must complete all of tasks up-to (and including) task 3a, during the regular lab. Tasks 3b and 4 may be done in a completion lab, if you run out of time.

Task 0

Task 0a:

Download the following file: [a.out](#). Answer the following questions (be prepared to explain your answers to the lab instructor):

1. Where is the entry point specified, and what is its value?
2. How many sections are there in a.out?
3. What is the size of the .text section?
4. Does the symbol `_start` occur in the file? If so, where is it mapped to in virtual memory?
5. Does the symbol `main` occur in the file? If so, where is it mapped to in virtual memory?
6. Where in the file does the code of function "main" start?

Task 0b

Write a program called *hexeditplus*:

```
./hexeditplus
```

The *hexeditplus* program performs operations (read and write) on files and memory. File operations are done on a file *filename* as defined below. Each operation is done in units of *size* bytes, which indicates a unit size, i.e. the number of bytes we want to use as the basic unit in each operation of our program, such as "display memory contents". Size can be either 1, 2 or 4, with 1 as the default.

For the ease of checking your code, your program will allocate a fixed size buffer of 4k bytes at the beginning of execution.

Whenever a memory address is set to 0 in the parameters of **hexeditplus** it will refer to the memory address of this buffer. The variables *size*, *filename* (initially null), and the memory buffer should all be global variables.

First, define a menu for the user with a number of predefined functions (as done in [Lab 2](#)), to which we will add functions as we go. The program prints the menu, obtains a choice from the user, acts on it, and repeats infinitely. For example, if functions: "Toggle Debug Mode", "Set File Name", "Mem Display", "Load Into Memory", "Save Into File" and "Quit" are available, then the command line:

```
./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Quit
```

For this part, use an array with the above menu names and pointers to appropriate functions that implement each option.

At this point implement "Toggle Debug Mode", "Set File Name", "Set Unit Size", and "Quit".

Toggle debug mode means turn the debug flag on (if it is currently off, which it is in the initial state), and print "Debug flag now on". If the debug flag is on, this function prints "Debug flag now off", and turns the flag off. When the debug mode is on, you should print the value of the global variables: unit size, file name, and buffer address, every time just before the menu is printed.

Optional (bonus): the debug toggle function can require entering a secret password of your choice, and only activates debug mode if the password is provided by the user. Such passwords are known in the computer gaming industry as "cheat codes" (so called by the users).

Set File Name queries the user for a file name, and stores it in a globally accessible buffer. You may assume that the file name is no longer than 100 characters. If debug mode is on, the function should also print: "Debug: file name set to 'filename' " (obviously, replacing 'filename' with the actual name).

The Set Unit Size option sets the size variable. The steps are:

1. Prompt the user for a number.
2. If the value is valid (1, 2, or 4), set the size variable accordingly.
3. If debug mode is on, print "Debug: set size to x", with x the appropriate size.
4. If not valid, print an error message and leave size unchanged.

Quit is a function that prints "quitting" (in debug mode), and calls `exit(0)` to quit the program.

The rest of the functions will be written in the next tasks. The menu should be extensible, you will change and extend it in each sub-task of task 1. It should be printed using a loop iterating over the menu array, and be `{NULL, NULL}` terminated.

All functions should be of the form:

```
void fun( );
```

Be sure to implement this code and test it carefully before the lab (that is why you have the debug option), as you will need to extend it during the lab!

Task 1: hexeditplus

In this task we will write our own version of *hexedit* for working with binary files. You will extend your code from task 0b.

Note: You should verify that there is no error when opening a file. In case of an error, you should print a message and abort the rest of the operation.

For this task you will be working with the following ELF file: [abc](#).

Task 1a: Mem Display

Write the function for the "Mem Display" option:

This option displays *length* units from **hexeditplus** program memory (**not** from the file *filename*), starting at virtual address *address*.

The steps are:

1. Prompt the user for an address (in hexadecimal) and a length (in decimal). If debug flag is on, it should also print "Default value: (the default buffer address)".
2. Display, in hexadecimal, *length* units from **hexeditplus** memory starting from *address*.

For example, the command line:

```
./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Mem Display
4-Quit
```

After the user chooses 3 with *address* "8048570" and *length* "5" - your program will print the 5 units, starting from the address 0x8048570 to 0x8048579.

The prompt should look as follows:

```
3
Please enter <address> <length>
8048570 5
```

The output should look something like this (for a unit size of 2):

```
ED31 895E 83E1 F0E4 5450
```

Note that if the user enters 0 as the address, your program should use the address of the default buffer, (say 0xffbf4054 if it is on the stack), instead.

To help debug your code, you can use the command:

```
readelf -S hexeditplus
```

To see where different parts of *hexeditplus* reside in the ELF executable file, and where will it reside in the memory when it is executed. You can then compare your output to the expected output.

Remember

- To read *address* and *length* use *fgets* and then *sscanf*, rather than *scanf* directly.
- *address* is entered in hexadecimal representation.

Task 1b: Load Into Memory

Write the function for the "Load Into Memory" option, which works as follows:

- Check if filename is null, and if it is print an error message and return.
- Open *source-file* for reading. If this fails, print an error message and return.
- Prompt the user for memory address *mem-address*, *location* (in hexadecimal) and *length* (in decimal).
- If debug flag is on, print the filename, as well as *mem-address*, *location*, and *length*.
- Loads into **hexeditplus** program memory at address *mem-address*, *length* units from *source-file* starting from position *location*.
- If *mem-address* = 0, the program will load the bytes into the default *fixed* buffer allocated by **hexeditplus** program.
- Close the file.
- Displays the address of the beginning of the buffer to which the units were loaded.

For example, the command:

```
./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
```

```

1-Set File Name
2-Set Unit Size
3-Mem Display
4-Load Into Memory
5-Quit

```

Assume that the user has already set the file name to "abc". If the user chooses 4, he is prompted for *mem-address*, *location* and *length*. It should look as follows:

```

4
Please enter <mem-address> <location> <length>
0 12F 10

```

The program should open the file [abc](#) and load the 10 units, from byte 303 to byte 312 in the file into memory. The output should look like:

```
Loaded 10 units into 0xffbf4054
```

You should use *hexedit* and task 1a (Mem Display), to verify that your code works correctly. Here is some of *hexedit*'s output for the file [abc](#), verify that you understand why the output is as it is.

00000070	01 00 00 00	01 00 00 00	00 00 00 00	00 80 04 08
00000080	00 80 04 08	EC 05 00 00	EC 05 00 00	05 00 00 00
00000090	00 10 00 00	01 00 00 00	14 0F 00 00	14 9F 04 08
000000A0	14 9F 04 08	0C 01 00 00	14 01 00 00	06 00 00 00
000000B0	00 10 00 00	02 00 00 00	28 0F 00 00	28 9F 04 08(...(...
000000C0	28 9F 04 08	C8 00 00 00	C8 00 00 00	06 00 00 00	(.....
000000D0	04 00 00 00	04 00 00 00	48 01 00 00	48 81 04 08H...H...
000000E0	48 81 04 08	44 00 00 00	44 00 00 00	04 00 00 00	H...D...D.....
000000F0	04 00 00 00	51 E5 74 64	00 00 00 00	00 00 00 00Q.td.....
00000100	00 00 00 00	00 00 00 00	00 00 00 00	06 00 00 00
00000110	04 00 00 00	52 E5 74 64	14 0F 00 00	14 9F 04 08R.td.....
00000120	14 9F 04 08	EC 00 00 00	EC 00 00 00	04 00 00 00
00000130	01 00 00 00	2F 6C 69 62	2F 6C 64 2D	6C 69 6E 75/lib/ld-linu
00000140	78 2E 73 6F	2E 32 00 00	04 00 00 00	10 00 00 00	x.so.2.....
00000150	01 00 00 00	47 4E 55 00	00 00 00 00	02 00 00 00GNU.....
00000160	06 00 00 00	0F 00 00 00	04 00 00 00	14 00 00 00
00000170	03 00 00 00	47 4E 55 00	C1 4E 4D 18	B9 A6 21 8FGNU..NM...!

Task 1c: Save Into File

Write the function for the "Save Into File" option, which works as follows:

This option replaces *length* * *size* bytes at *target-location* of *filename* with bytes from the **hexeditplus** memory address starting at virtual address *address*.

For example, the command:

```
./hexeditplus
```

Will print:

```

Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Mem Display
4-Load Into Memory
5-Save Into File
6-Quit

```

When the user chooses option 5, the program should query the user for:

- *s-address* (source memory address, in hexadecimal), *s-address* can be set to 0, in which case, the source address is that of the fixed buffer.
- *t-location* (target file offset, in hexadecimal),
- *length* (number of *size* * bytes, in decimal).

Implement the checks that the file can be opened (for writing and NOT truncating), and print appropriate debug messages in debug mode as in the previous task. Close the file after writing.

For example, after the file name was set to "abc", choosing option "5-Save Into File" using *s-address* 960c170, *t-location* 33 and *length* 4, the program should read *length* = 4 bytes from memory, starting at address 0x960c170, and write them to the file *abc*, starting from offset 0x33 (overwriting what was originally there). It should look as follows:

```
5
Please enter <s-address> <t-location> <length>
960c170 33 4
```

Note that the target file is specified by global file name variable.

Also observe that after you execute this option, **only** *length* * *size* bytes of the file *targetfile* should be changed.

If <t-location> is greater than the size of <target-file> you should print an error message and not copy anything.

Task 1d: Mem Modify

Write the function for the "Mem Modify" option:

This option replaces a unit at *address* in **hexeditplus** memory by *val*.

The steps are:

1. Prompt the user for *address* and *val* (all in hexadecimal).
2. If debug mode is on, print the address and val given by the user.
3. Replace a unit at *address* in the memory with the values given by *val*.

For example, the command:

```
./hexeditplus
```

Will print:

```
Choose action:
0-Toggle Debug Mode
1-Set File Name
2-Set Unit Size
3-Mem Display
4-Load Into Memory
5-Save Into File
6-Mem Modify
7-Quit
```

When the user chooses option 6, the program should query the user for:

- *address* (memory address, in hexadecimal), *address* can be 0, in which case, the address used is that of the fixed buffer.

- *val* (new value, in hexadecimal)

For example, if unit size was set to 4, choosing option "6-Mem Modify" using *address* 960c170, *val* 804808a, will overwrite the 4 bytes starting at location 960c170, with the new value 804808a. It should look as follows:

```
6
Please enter <address><val>
960c170 804808a
```

You can test the correctness of your code using task1a - "Mem Display".

Task 2: Reading ELF

Task 2a

Download the following file: [chezi](#).

chezi is an executable ELF file. It does not run as expected. Your task is to understand the reason for that.

Do the following:

1. Run the file.
2. Which function precedes main in execution ? (hint: The [assembly code in Lab 4](#)).
3. What is the virtual address to which this function is loaded (hint: use `readelf -s`)

Task 2b

Use your *hexeditplus* program from task 1 to display the entry point of a file.

What are the values of *location/length*? How do you know that?

Use the edit functions from *hexeditplus* program to fix the *chezi* file, so that it behaves as expected.

You would need a combination of several options in *hexeditplu* to complete this task.

Task 3: Delving Deeper into the ELF Structure

task 3a

The goal of this task is to display the compiled code (in bytes) of the function *main*, in the *abc* executable above.

In order to do that, you need to:

1. find the offset (file location) of the function *main*.
2. find the size of the function *main*.
3. use your *hexeditplus* program to display the content of that function on the screen.

Finding the needed information:

1. Find the entry for the function `main` in the symbol table of the ELF executable (`readelf -s`).
2. In that reference you will find both the size of the function and the function's virtual address and section number.
3. In the section table of the executable, find the entry for the function's section (`readelf -S`).
4. Find both the section's virtual address (Addr), and the section's file offset (Off).
5. Use the above information to find the file offset of the function.

Task 3b-asm

This task is only for students of the architecture and splab course.

What are the first two machine instructions in function `main`, stated in assembly language? I.e. you need to manually dis-assemble these first two instructions.

You can use the opcode information in the [nasm manual](#) .

Have a look at [Practical session 3](#) page 9, before you delve into nasm manual.

Task 3b-splab

This task is for students registered for splab course only and NOT architecture.

Hack this executable file so that it does nothing when it is run: replace the code of the `main` function by [NOP](#) instructions.

Make sure you do NOT override the `ret` instruction ([Opcode](#) : c3) in `main`.

Alternately, you can plant just one `ret` instruction (where?).

Task 4: Hacking: installing a patch using hexeditplus

The following file [ntsc](#) was meant to be a digit counter. Download it, and run it in the command-line.

```
./ntsc aabbaba123baacca  
./ntsc 1112111
```

What is the problem with the file? (hint, try this string: 0123456789)

Create a new program with a correct digit counter function (should get a `char*` and return an `int`), compile and test it. (remember to compile with the `-m32` flag in order to produce an ELF compatible with 32bits).

Use *hexeditplus* to replace (patch) the buggy `digit_cnt` function in the `ntsc` file with the corrected version from the new program.

You should do it using options 4 & 5 in *hexeditplus*.

(think: are there any kinds of restrictions on the code you wrote for the `digit_cnt` function?)

Explain how you did it, and show that it works.

Deliverables:

Tasks 1,2, and 3a must be completed during the regular lab. Tasks 3b and 4 may be done in a completion lab, but only if you run out of time during the regular lab. The deliverables must be submitted until the end of the day.

